

HIVE-13680: Provide a way to compress ResultSets

Kevin Liew

May 13, 2016

Abstract

Hive data pipelines invariably involve JDBC/ODBC drivers which access Hive through its Thrift interface. Consequently, any enhancement to the Hive's Thrift data pipeline will benefit all users.

Prior to HIVE-12049, HiveServer2 would read a full ResultSet from HDFS before deserializing and re-serializing into Thrift objects for RPC transfer. Following our enhancement, task nodes serialize Thrift objects from their own block of a ResultSet file. HiveServer2 simply transfers these objects to a remote client. This parallel serialization strategy reduced latency in the data pipeline.

However, network load is often the most scarce resource in a system. As a further enhancement, this load can be eased by having task nodes compress their own block of a ResultSet as part of the serialization process.

1 Introduction

The changes proposed herein draw from Rohit Dholakia's design document and patches for HIVE-10438, which implemented compression on HiveServer2 rather than in the task nodes. Now that HIVE-12049 has been committed, compression can take place in parallel on the task nodes.

Our goals for this enhancement are to:

- improve performance out-of-the-box for new clients
- maintain compatibility with old clients
- provide flexibility yet security
- confer a simple interface

2 Design Overview

2.1 Compressor Interface

We will define a compressor interface which must be implemented by compressor plugins. A Snappy compressor will compress all data-types using

the Snappy algorithm. Type-specific compression can be achieved by implementing a compressor that delegates compression to other compressors based on a case-switch block.

2.2 Client-Server Negotiation

The client will specify a preferred compressor upon connection. If the server does not have that compressor plugin, it will request the jar from the client. The client should send the compressor plugin and all dependencies.

To prevent breaking compatibility with old clients, the server will not compress results unless the client has requested compression. However, new versions of beeline will try to negotiate a compressor scheme by default.

2.3 Server Security

We place the onus of security on the server administrator, recommending Kerberos authentication with user-impersonation to limit the jar file to the same permissions as the authenticated user. The server administrator also has the option to disable compressor plugins altogether.

2.4 Configuration options

Option	Default	Description
hive.resultSet.compressor.enabled	true	Enable or disable compressor negotiation

Table 1: Configuration options

3 Implementation

3.1 Compressed RowSet Structure

Compressors will operate on row-batches in ThriftJDBCBinarySerDe. To model a compressed column, we proposed CompressedColumnBasedSet.

In the event that a column is not compressible, the column will be a normal ColumnBasedSet.

3.1.1 HiveServer2 Class

TBD

3.1.2 Thrift Object

TBD

3.2 Compressor Interface

TBD

4 Custom Compressors

A default compressor will be implemented for Snappy. To use other compression algorithms or to have type-specific compression, the user must implement the compressor interface.

more TBD